

Lecture12: Value Estimation

Notes taken by [squarezhong](#)

Repo address: [squarezhong/SDM5008-Lecture-Notes](#)

Lecture12: Value Estimation

Algorithm: MC estimation, for estimating state values

Notation

Process

Examples

Incremental Implementation

From Monte-Carlo (MC) to Temporal Difference (TD)

TD with n -step returns

Rules for action-values

Trade-off between bias and variance:

λ -return: TD(λ) algorithm

Definition

Notation

Derivation

Discussion

Finite Case

Summary

Further Expansion

Value Parameterization

Truncated Policy Evaluation

Difference between leveraging state-values and action-values

Algorithm: MC estimation, for estimating state values

Notation

For a finite trajectory $\tau = \{s_t, a_t, r_{t+1}\}_{t=0}^{T-1}$

$$G_t = G(s_t) = R(\tau|S_0 = s_t) = r_{t+1} + \gamma r_{t+2} + \cdots + \gamma^{T-1} r_T$$

Corollary:

$$G(s_t) = r_{t+1} + \gamma G(s_{t+1})$$

Process

- Input: policy π to be evaluate
- Output: $V_{\pi}^*(s)$
- Init:
 - $V(s) \in \mathbb{R}, \forall s \in \mathbb{S}$
 - an empty list $R(s), \forall s \in \mathbb{S}$

Loop for each episode: (termination: $\max V_{\pi}^{[k+1]}(s_t) - V_{\pi}^{[k]}(s_t) \leq \Delta$)

generate $\tau = \{s_t, a_t, r_{t+1}\}_{t=0}^{T-1}$ using π

$G_t \leftarrow 0$

Loop for each time step $t = T - 1, T - 2, \dots, 0$:

$G_t \leftarrow r_{t+1} + \gamma G_{t+1}$

$R(s_t)$ append G_t

$\hat{V}_\pi(s_t) \leftarrow \text{average}(R(s_t))$

Above is the **"every visit"** condition.

This calculates the value function for all possible states?

Examples

- [Frozen Lake](#)
- [Example Code](#)

Incremental Implementation

$$V_\pi^{[m+1]}(s_t) = V_\pi^{[m]} + \frac{1}{m+1} \left(G_t^{[m+1]} - V_\pi^{[m]}(s_t) \right)$$

New Estimation \leftarrow Old Estimation + α (New Observation - Old Estimation)

This significantly reduces the storage requirements.

Derivation:

$$\begin{aligned} V_\pi^{[m+1]}(s_t) &= \frac{1}{m+1} \sum_{i=1}^{m+1} G_t^{[i]} \\ &= \frac{1}{m+1} \left(\sum_{i=1}^m G_t^{[i]} + G_t^{[m+1]} \right) \\ &= \frac{1}{m+1} \cdot m \cdot \frac{1}{m} \left(\sum_{i=1}^m G_t^{[i]} + G_t^{[m+1]} \right) \\ &= \frac{m}{m+1} V_\pi^{[m]}(s_t) + \frac{1}{m} G_t^{[m+1]} \\ &= \frac{1}{m+1} \left((m+1) V_\pi^{[m]}(s_t) - V_\pi^{[m]}(s_t) + G_t^{[m+1]} \right) \\ &= V_\pi^{[m]}(s_t) + \frac{1}{m+1} \left(G_t^{[m+1]} - V_\pi^{[m]}(s_t) \right) \end{aligned}$$

From Monte-Carlo (MC) to Temporal Difference (TD)

- Limitation of MC methods: require finished episodes
- TD: $G_t \approx r_{t+1} + \gamma \hat{V}_\pi(s_{t+1})$, applies if some applications have very long episodes

$$V_\pi^{[m+1]}(s_t) = V_\pi^{[m]}(s_t) + \alpha \left(\underbrace{r_{t+1} + \gamma \hat{V}_\pi(s_{t+1})}_{\text{TD Target}} - V_\pi^{[m]}(s_t) \right)$$

The right side of α is call **TD Error**

TD with n -step returns

- 1-step TD: $G_t \approx r_{t+1} + \gamma \hat{V}_\pi(s_{t+1})$
- 2-step TD: $G_t \approx r_{t+1} + \gamma r_{t+2} + \gamma^2 \hat{V}_\pi(s_{t+2})$
- n -step TD: $G_t \approx r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{n-1} r_{t+n} + \gamma^n \hat{V}_\pi(s_{t+n})$
- Monte Carlo methods could be considered as ∞ -step TD methods.

Rules for action-values

- MC: $Q_\pi^{[m+1]}(s_t, a_t) = \frac{1}{m+1} \sum_{i=1}^{m+1} G_t^{[i]}$
- MC-incremental: $Q_\pi^{[m+1]}(s_t, a_t) = Q_\pi^{[m]}(s_t, a_t) + \frac{1}{m+1} (A_t^{[m+1]} - Q_\pi^{[m]}(s_t, a_t))$
- n -step TD: $G_t \approx r_{t+1} + \dots + \gamma^{n-1} r_{t+n} + \gamma^n \hat{Q}_\pi(s_{t+n}, a_{t+n})$

Trade-off between bias and variance:

- MC methods: unbiased, high variance
- TD methods: biased, low variance
- n -step TD methods: the choice of n serves as a **trade-off** between bias and variance in value estimation

λ -return: TD(λ) algorithm

Definition

Infinite Case:

$$G_t^\lambda = (1 - \lambda) \sum_{l=1}^{\infty} \lambda^{l-1} G_{t:t+l}$$

Estimating the value functions using the TD method computed with the λ -return.

Another method to trade off between the bias and variance of the estimator. Empirically better than fixed n -step return.

Notation

$G_{t:t+n}$ is equivalent to the n -step TD

$$G_{t:t+n} = r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{n-1} r_{t+n} + \gamma^n \hat{V}_\pi(s_{t+n})$$

Derivation

$$\begin{aligned} G_t^\lambda &= (1 - \lambda)(G_{t:t+1} + \lambda G_{t:t+2} + \lambda^2 G_{t:t+3} + \dots) \\ &= (1 - \lambda)(r_{t+1} + \gamma \hat{V}_\pi(s_{t+1})) + (1 - \lambda)\lambda(r_{t+1} + \gamma r_{t+2} + \gamma^2 \hat{V}_\pi(s_{t+2})) + \dots \\ &= (1 - \lambda)[(1 + \lambda + \lambda^2 + \dots)r_{t+1} + \gamma\lambda(1 + \lambda + \lambda^2) + \gamma^2\lambda^2(1 + \lambda + \lambda^2)] \\ &\quad + (1 - \lambda)(\gamma \hat{V}_\pi(s_{t+1}) + \gamma^2\lambda \hat{V}_\pi(s_{t+2}) + \gamma^3\lambda^2 \hat{V}_\pi(s_{t+3})) \\ &= \sum_{l=1}^{\infty} \left[(\gamma\lambda)^{l-1} r_{t+l} + (1 - \lambda)\gamma^l \lambda^{l-1} \hat{V}_\pi(s_{t+l}) \right] \end{aligned}$$

Intermediate expressions used:

- $1 + \lambda + \lambda^2 + \dots = \lim_{n \rightarrow \infty} \frac{1 - \lambda^{n+1}}{1 - \lambda} = \frac{1}{1 - \lambda}$

Discussion

Observe the form

$$G_t^\lambda = \sum_{l=1}^{\infty} \left[(\gamma\lambda)^{l-1} r_{t+l} + (1-\lambda)\gamma^l \lambda^{l-1} \hat{V}_\pi(s_{t+l}) \right]$$

2 Special cases:

- $\lambda = 0$: $G_t^0 = r_{t+1} + \gamma \hat{V}_\pi(s_{t+1})$
This is 1-step return TD
- $\lambda = 1$: $G_t^1 = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots$
This is Monte Carlo (∞ -step return)

When $0 < \lambda < 1$:

- λ -return offers a compromise between bias and variance controlled by λ

Finite Case

assume all rewards after step T are 0

$$G_{t:t+n} = G_{t:t+T}, \forall n \geq T-t$$

Then the general form of λ -return under finite case is:

$$\begin{aligned} G_t^\lambda &= (1-\lambda) \sum_{l=1}^{T-t-1} \lambda^{l-1} G_{t:t+l} + (1-\lambda) \sum_{l=T-t}^{\infty} (1+\lambda+\lambda^2+\dots) \lambda^{l-1} G_{t:T} \\ &= \sum_{l=1}^{T-t-1} \lambda^{l-1} G_{t:t+l} + \lambda^{T-t-1} G_{t:T} \end{aligned}$$

Summary

$$\hat{V}_\pi^{[m+1]}(s_t) \leftarrow \hat{V}_\pi^{[m]}(s_t) + \frac{1}{m+1} (G_t^{[m+1]} - \hat{V}_\pi^{[m]}(s_t))$$

To estimate $G_t^{[m+1]}$, we can use:

- ∞ -step return \rightarrow MC
- (n) 1-step return \rightarrow TD
- λ -return \rightarrow TD(λ)

Further Expansion

Value Parameterization

Q: If there are a large number of or infinitely many states, it is difficult or impossible to storage all state values.

Solution: Using a **function approximator** (e.g. neural network) to estimate the state-value function $\hat{V}_\pi(s; w)$


Find w^* s.t. $w^* = \operatorname{argmin} E_\pi[(y_t - \hat{V}_\pi(s_t; w))^2]$

- y_t is the true value but may be not accessible
- we can use TD-target $r_{t+1} + \gamma \hat{V}_\pi(s_{t+1}; w)$ to replace y_t

MSE Loss:

$$L(w) = \frac{1}{N} \sum_i (r_{t+1}^{[i]} + \gamma \hat{V}_\pi(s_{t+1}^{[i]}; w) - \hat{V}_\pi(s_t^{[i]}; w))$$

Truncated Policy Evaluation



Truncate early before evaluation converges to improve the policy improvement speed.

Difference between leveraging state-values and action-values

Answers from ChatGPT 4o

In practice, **state-values** are used when you're interested in evaluating or improving a policy in a model-based way, whereas **action-values** are more commonly used in model-free reinforcement learning algorithms, allowing agents to directly learn the optimal policy by interacting with the environment.