

# Lecture14: Proximal Policy Optimization (PPO)

Notes taken by [squarezhong](#)

Repo address: [squarezhong/SDM5008-Lecture-Notes](#)

## Lecture14: Proximal Policy Optimization (PPO)

Review: Policy Optimization

Review: Vanilla Policy Gradient (VPG)

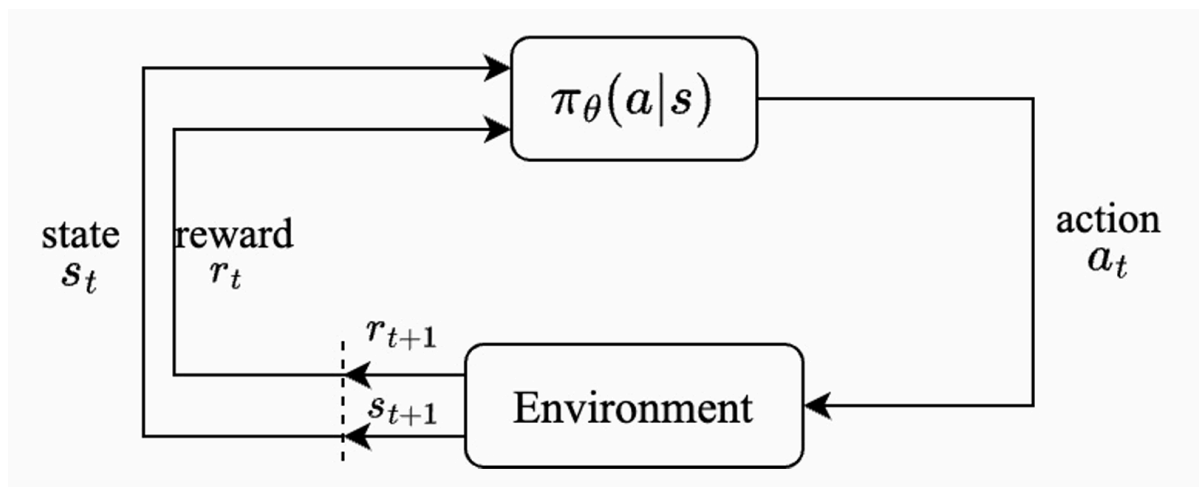
Importance Sampling and Surrogate Loss

Proximal Policy Optimization (PPO)

General Advantage Estimation (GAE)

Code Example

## Review: Policy Optimization



- $\pi(a|s)$ : Probability of action  $a$  in state  $s$
- Objective: Maximize expected cumulative reward

$$\max_{\theta} \mathbf{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^T \gamma^t r_t \right]$$

## Review: Vanilla Policy Gradient (VPG)

**Input:** Initial  $\theta_0, V_{\phi}$

**Repeat:**

1. **Collect trajectories:**  $\tau \sim \pi_{\theta}$
2. **For each timestep  $t$ :**
  - $G_t = \sum_{t'=t}^{T-1} \gamma^{t'-t} r_{t'}$  (return)
  - $A_t = G_t - V_{\phi}(s_t)$  (advantage function)
3. **Update  $V_{\phi}$ :**

$$V_{\phi} = \operatorname{argmin}_{\phi} \sum \|V_{\phi}(s_t) - G_t\|^2$$

4. **Policy Update:**

$$\theta = \theta + \alpha \sum \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) A_t$$

## Importance Sampling and Surrogate Loss

- Policy Gradient

$$J(\theta) = \mathbf{E}[\log \pi_{\theta}(a_t | s_t) \cdot A(s_t, a_t)]$$

$$\nabla_{\theta} J(\theta) = \mathbf{E}[\nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot A(s_t, a_t)]$$

- Surrogate Loss Function (替代损失函数):

$$L_{\text{surr}} = \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\text{old}}(a_t | s_t)} A(s_t, a_t)$$

Here we prove that  $L_{\text{surr}}$  can replace policy gradient to update the policy

Prove that  $\nabla_{\theta} \mathbf{E}(L_{\text{surr}}) = \nabla_{\theta} J(\theta)$  when  $\theta = \theta_{\text{old}}$

$$\begin{aligned} \nabla_{\theta} \mathbf{E}(L_{\text{surr}}) &= \nabla_{\theta} \mathbf{E} \left[ \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\text{old}}(a_t | s_t)} \cdot A(s_t, a_t) \right] \\ &= \mathbf{E} \left[ \nabla_{\theta} \pi_{\theta}(a_t | s_t) \cdot \frac{1}{\pi_{\text{old}}(a_t | s_t)} \cdot A(s_t, a_t) \right] \\ &= \mathbf{E} \left[ \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\text{old}}(a_t | s_t)} \cdot \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot A(s_t, a_t) \right] \\ &\approx \mathbf{E} [\nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot A(s_t, a_t)] \\ &= \nabla_{\theta} J(\theta) \end{aligned}$$

- It update policy via:

$$\theta_{k+1} = \operatorname{argmax}_{\theta, a \sim \pi_{\theta_k}} \mathbf{E} [L(s, a, \theta_k, \theta)]$$

- Typically taking multiple steps of SGD to maximize the objective.

## Proximal Policy Optimization (PPO)

- PPO introduces a clipping mechanism to prevent large, destabilizing updates.
- Clipped objective:

$$L(s, a, \theta_k, \theta) = \min \left( \frac{\pi_{\theta}(a | s)}{\pi_{\theta_k}(a | s)} A^{\pi_{\theta_k}}(s, a), \operatorname{clip} \left( \frac{\pi_{\theta}(a | s)}{\pi_{\theta_k}(a | s)}, 1 - \epsilon, 1 + \epsilon \right) A^{\pi_{\theta_k}}(s, a) \right)$$

- $\epsilon$  is hyperparameter which roughly says how far away the new policy is allowed to go from the old.

- Clipped objective in another form:

$$L(s, a, \theta_k, \theta) = \min \left( \frac{\pi_{\theta}(a | s)}{\pi_{\theta_k}(a | s)} A^{\pi_{\theta_k}}(s, a), g(\epsilon, A^{\pi_{\theta_k}}(s, a)) \right),$$

$$\text{where } g(\epsilon, A) = \begin{cases} (1 + \epsilon)A & \text{if } A \geq 0, \\ (1 - \epsilon)A & \text{if } A < 0. \end{cases}$$

- When the advantage is positive:

- Action  $a$  becomes more likely but has a limit.
- Good action appears more but not too more
- When the advantage is negative:
  - Action  $a$  becomes less likely but has a limit.
  - Bad action appears less but not too less.

Pseudo code:

---

### Algorithm 1 PPO-Clip

---

**Require:** Initial policy parameters  $\theta_0$ , initial value function parameters  $\phi_0$

- 1: **for**  $k = 0, 1, 2, \dots$  **do**
- 2:   Collect set of trajectories  $\mathcal{D}_k = \{\tau_i\}$  by running policy  $\pi_k = \pi(\theta_k)$  in the environment.
- 3:   Compute rewards-to-go  $\hat{G}_t$ .
- 4:   Compute advantage estimates,  $\hat{A}_t$  (using any method of advantage estimation) based on the current value function  $V_{\phi_k}$ .
- 5:   Update the policy by maximizing the PPO-Clip objective:

$$\theta_{k+1} = \arg \max_{\theta} \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \min \left( \frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_k}(a_t|s_t)} A^{\pi_{\theta_k}}(s_t, a_t), g(\epsilon, A^{\pi_{\theta_k}}(s_t, a_t)) \right),$$

typically via stochastic gradient ascent with Adam.

- 6:   Fit value function by regression on mean-squared error:

$$\phi_{k+1} = \arg \min_{\phi} \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \left( V_{\phi}(s_t) - \hat{G}_t \right)^2,$$

typically via some gradient descent algorithm.

- 7: **end for**
- 

## General Advantage Estimation (GAE)

Advantage  $\hat{A}_t^{(k)}$ :

$$\hat{A}_t^{(k)} = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots + \gamma^k V(s_{t+k}) - V(s_t)$$

**GAE** takes a weighted average of  $\hat{A}_t^{(k)}$  to balance bias and variance.

$$\hat{A}_t^{(k)} = A_t^{\text{GAE}} = \sum_k w_k \hat{A}_t^{(k)}$$


---

The derivation process of a recursive form

- $w_k = \lambda^{k-1}$ ,  $\lambda \in [0, 1]$
- $\delta_t = r_t + \gamma V(s_{t+1}) - V(s_t)$

Then:

- $k = 1$ :  $\hat{A}_t^{(1)} = \delta_t$
- $k = 2$ :  $\hat{A}_t^{(2)} = \delta_t + \gamma \delta_{t+1}$

- $k = 3$ :  $\hat{A}_t^{(3)} = \delta_t + \gamma\delta_{t+1} + \gamma^2\delta_{t+2}$

Takes a weighted average:

$$\begin{aligned}
 \hat{A}_t &= (1 - \lambda)(\hat{A}_t^{(1)} + \lambda\hat{A}_t^{(2)} + \lambda^2\hat{A}_t^{(3)} + \dots) \\
 &= (1 - \lambda)(\delta_t + \lambda\delta_t + \gamma\lambda\delta_{t+1} + \lambda^2\delta_t + \gamma\lambda^2\delta_{t+1} + \gamma^2\lambda^2\delta_{t+2} + \dots) \\
 &= (1 - \lambda)[\delta_t(1 + \lambda + \lambda^2 + \dots) + \gamma\delta_{t+1}(\lambda + \lambda^2 + \dots) + \gamma^2\delta_{t+2}(\lambda^2 + \lambda^3 + \dots) + \dots] \\
 &= \delta_t + \gamma\lambda\delta_{t+1} + (\gamma\lambda)^2\delta_{t+2} + \dots \\
 &= \delta_t + \gamma\lambda\hat{A}_{t+1}
 \end{aligned}$$

Special Case:

- $\lambda = 0$ :  $\hat{A}_t = \delta_t = \hat{A}_t^{(1)}$
- $\lambda = 1$ :  $\hat{A}_t = \delta_t + \gamma\hat{A}_{t+1} = \hat{A}_t^{(\infty)}$

---

## Code Example

- [Code Example](#)